

# performance dataset

By Hatem Elgenedy

December 7, 2025

```
[ ]: # import necessary libraries
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
import pandas as pd
from sklearn.model_selection import train_test_split as split
from sklearn.preprocessing import StandardScaler
```

```
[ ]: # load the dataset
df = pd.read_csv("/content/employee_performance_pro.csv")
df.head(5)
```

```
[ ]: EmployeeID          Name  Gender  Age Department      JobRole \
0           1    Steven Barnett   Other   57   Finance       Auditor
1           2  Christopher Benson Female   26     Sales  Sales Executive
2           3      Norman Lane   Other   59     Support  Helpdesk
3           4      Rita Walker Female   43        HR  HR Executive
4           5      Judith Ware   Male   52     Sales Account Manager

      EducationLevel  JoiningDate  CountryCode  Country ... LastLeaveDate \
0                 2  2016-05-05         91    India ... 2024-07-03
1                 2  2014-08-20          1  Canada ... 2024-01-05
2                 1  2010-05-17          1  Canada ... 2024-11-27
3                 3  2015-06-20         49  Germany ... 2024-07-15
4                 3  2019-08-20         49  Germany ... 2024-12-17

      LeaveDayName  ProjectsHandled  TrainingHours CustomerSatisfaction \
0      Wednesday                  13             16                NaN
1      Friday                   15             44              8.0
2      Wednesday                  7              62              7.0
3      Monday                   15              8                NaN
4      Tuesday                  10              57              8.0

      LastPromotionYear  YearsAtCompany  WorkLifeBalanceScore  PerformanceRating \
0            2020                  9             1.90                  4
1            2014                  11             5.03                  1
2            2012                  15             4.83                  4
```

```
3          2016      10      4.60      4
4          2022       6      3.73      4
```

```
AttritionRisk
0          No
1         Yes
2          No
3          No
4          No
```

[5 rows x 24 columns]

```
[ ]: df.info() # check for null values and data types
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 24 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   EmployeeID      500 non-null    int64  
 1   Name             500 non-null    object  
 2   Gender           500 non-null    object  
 3   Age              500 non-null    int64  
 4   Department       500 non-null    object  
 5   JobRole          500 non-null    object  
 6   EducationLevel   500 non-null    int64  
 7   JoiningDate     500 non-null    object  
 8   CountryCode     500 non-null    int64  
 9   Country          500 non-null    object  
 10  PhoneNumber      500 non-null    int64  
 11  MonthlySalary   500 non-null    int64  
 12  OvertimeHoursPerMonth 500 non-null    int64  
 13  LeavesTaken     500 non-null    int64  
 14  LastLeaveDate   500 non-null    object  
 15  LeaveDayName    500 non-null    object  
 16  ProjectsHandled 500 non-null    int64  
 17  TrainingHours    500 non-null    int64  
 18  CustomerSatisfaction 181 non-null    float64 
 19  LastPromotionYear 500 non-null    int64  
 20  YearsAtCompany   500 non-null    int64  
 21  WorkLifeBalanceScore 500 non-null    float64 
 22  PerformanceRating 500 non-null    int64  
 23  AttritionRisk    500 non-null    object  
dtypes: float64(2), int64(13), object(9)
memory usage: 93.9+ KB
```

```
[ ]: df.shape # check the shape of the dataset
```

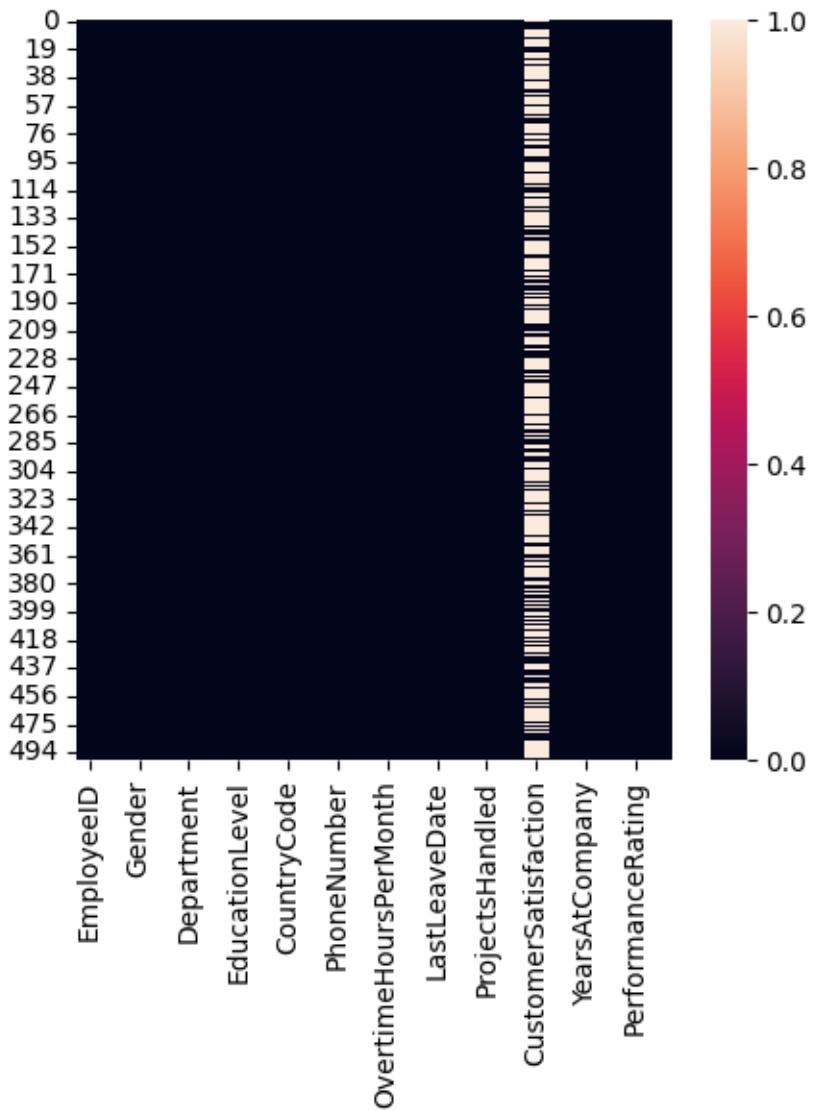
```
[ ]: (500, 24)
```

```
[ ]: df.isnull().sum() # check for null values in each column
```

```
[ ]: EmployeeID          0
Name                0
Gender              0
Age                 0
Department         0
JobRole             0
EducationLevel     0
JoiningDate        0
CountryCode         0
Country             0
PhoneNumber         0
MonthlySalary       0
OvertimeHoursPerMonth 0
LeavesTaken         0
LastLeaveDate       0
LeaveDayName        0
ProjectsHandled    0
TrainingHours       0
CustomerSatisfaction 319
LastPromotionYear   0
YearsAtCompany      0
WorkLifeBalanceScore 0
PerformanceRating   0
AttritionRisk       0
dtype: int64
```

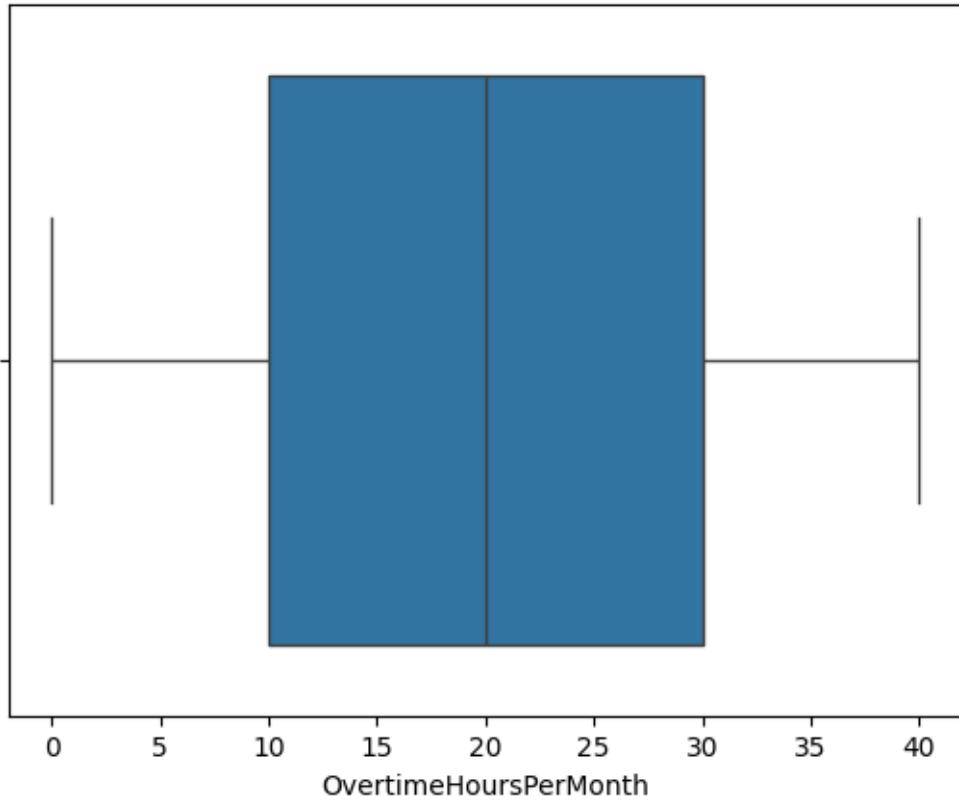
```
[ ]: plt.figure(figsize = (5,5)) # visualize null values using heatmap
sns.heatmap(df.isnull()) # heatmap for null values
```

```
[ ]: <Axes: >
```



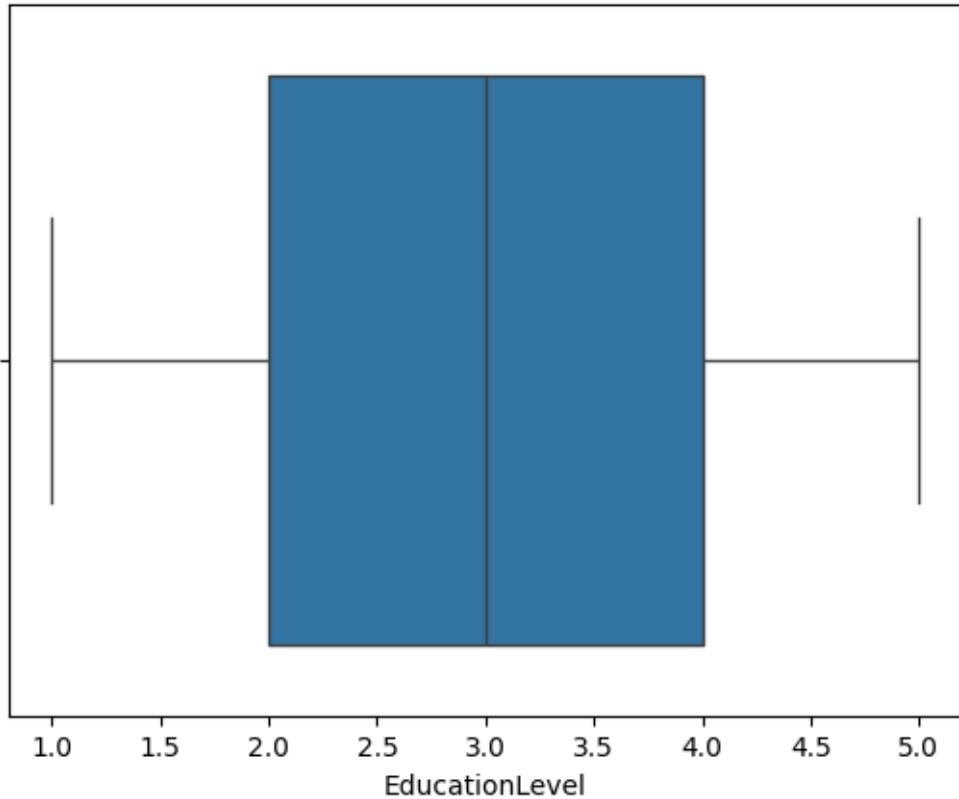
```
[ ]: df['OvertimeHoursPerMonth'] = pd.to_numeric(df['OvertimeHoursPerMonth'], errors='coerce') # convert to numeric, coercing errors to NaN
      sns.boxplot(x = 'OvertimeHoursPerMonth' , data = df) # boxplot to visualize outliers in OvertimeHoursPerMonth
```

```
[ ]: <Axes: xlabel='OvertimeHoursPerMonth'>
```



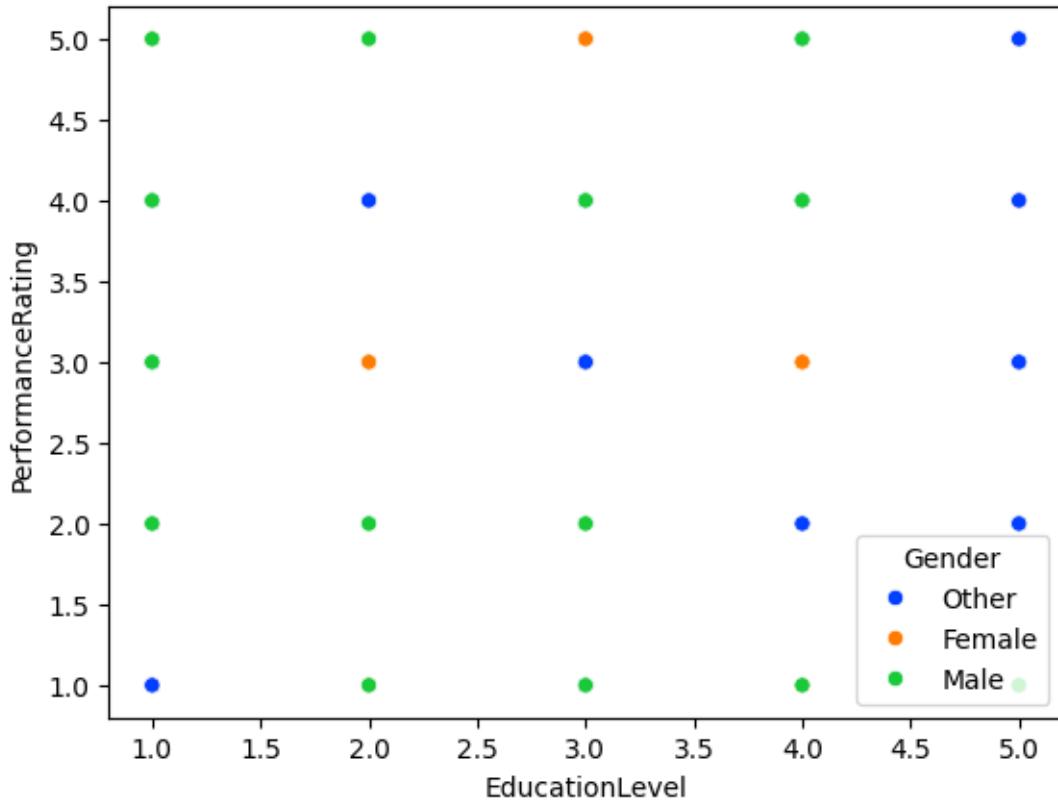
```
[ ]: df['EducationLevel'] = pd.to_numeric(df['EducationLevel'], errors = 'coerce') #  
    ↪convert EducationLevel to numeric  
sns.boxplot(x = 'EducationLevel' , data = df) # boxplot to visualize outliers  
    ↪in EducationLevel
```

```
[ ]: <Axes: xlabel='EducationLevel'>
```



```
[ ]: # scatter plot to visualize relationship between EducationLevel and PerformanceRating
      sns.scatterplot(data=df , x = 'EducationLevel' , y = 'PerformanceRating' , hue= 'Gender' , palette ='bright')
```

```
[ ]: <Axes: xlabel='EducationLevel', ylabel='PerformanceRating'>
```



```
[ ]: # Identify outliers in OvertimeHoursPerMonth using IQR method
df['OvertimeHoursPerMonth'] = pd.to_numeric(df['OvertimeHoursPerMonth'], errors='coerce')
data = df.dropna(subset = ["OvertimeHoursPerMonth"])
Q1 = np.percentile(data["OvertimeHoursPerMonth"] , 25)
Q3 = np.percentile(data["OvertimeHoursPerMonth"] , 75)
IQR = Q3 - Q1
# Calculate bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
# Filter out the outliers
outliers = data[(df["OvertimeHoursPerMonth"]< lower_bound) | (data["OvertimeHoursPerMonth"] > upper_bound)]

print("outliers :\n" , outliers)

outliers :
Empty DataFrame
Columns: [EmployeeID, Name, Gender, Age, Department, JobRole, EducationLevel,
JoiningDate, CountryCode, Country, PhoneNumber, MonthlySalary,
OvertimeHoursPerMonth, LeavesTaken, LastLeaveDate, LeaveDayName,
```

```
ProjectsHandled, TrainingHours, CustomerSatisfaction, LastPromotionYear,  
YearsAtCompany, WorkLifeBalanceScore, PerformanceRating, AttritionRisk]  
Index: []
```

```
[0 rows x 24 columns]
```

```
[ ]: df.shape # check the shape of the dataset
```

```
[ ]: (500, 24)
```

```
[ ]: df.info() # check for null values and data types
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 500 entries, 0 to 499  
Data columns (total 24 columns):  
 #   Column           Non-Null Count  Dtype     
---  --     
 0   EmployeeID      500 non-null    int64    
 1   Name             500 non-null    object    
 2   Gender           500 non-null    object    
 3   Age              500 non-null    int64    
 4   Department       500 non-null    object    
 5   JobRole          500 non-null    object    
 6   EducationLevel   500 non-null    int64    
 7   JoiningDate     500 non-null    object    
 8   CountryCode      500 non-null    int64    
 9   Country          500 non-null    object    
 10  PhoneNumber      500 non-null    int64    
 11  MonthlySalary    500 non-null    int64    
 12  OvertimeHoursPerMonth  500 non-null    int64    
 13  LeavesTaken     500 non-null    int64    
 14  LastLeaveDate    500 non-null    object    
 15  LeaveDayName    500 non-null    object    
 16  ProjectsHandled  500 non-null    int64    
 17  TrainingHours    500 non-null    int64    
 18  CustomerSatisfaction  181 non-null    float64    
 19  LastPromotionYear  500 non-null    int64    
 20  YearsAtCompany   500 non-null    int64    
 21  WorkLifeBalanceScore  500 non-null    float64    
 22  PerformanceRating  500 non-null    int64    
 23  AttritionRisk    500 non-null    object    
dtypes: float64(2), int64(13), object(9)  
memory usage: 93.9+ KB
```

```
[ ]: X = df.drop(columns = ['OvertimeHoursPerMonth']) # features  
y = df[['OvertimeHoursPerMonth']] # target variable
```

```
[ ]: X.columns # check feature columns
```

```
[ ]: Index(['EmployeeID', 'Name', 'Gender', 'Age', 'Department', 'JobRole',
   'EducationLevel', 'JoiningDate', 'CountryCode', 'Country',
   'PhoneNumber', 'MonthlySalary', 'LeavesTaken', 'LastLeaveDate',
   'LeaveDayName', 'ProjectsHandled', 'TrainingHours',
   'CustomerSatisfaction', 'LastPromotionYear', 'YearsAtCompany',
   'WorkLifeBalanceScore', 'PerformanceRating', 'AttritionRisk'],
  dtype='object')

[ ]: y.columns # check target variable column
```

```
[ ]: Index(['OvertimeHoursPerMonth'], dtype='object')
```

```
[ ]: y.head(2) # display first 2 rows of target variable
```

```
[ ]: OvertimeHoursPerMonth
0 33
1 24
```

```
[ ]: # split the dataset into training and testing sets
X_train , X_test , y_train , y_test = split(X, y , train_size = 0.8 ,random_state = 42)
X_train.shape
```

```
[ ]: (400, 23)
```

```
[ ]: # check the shape of the target variable training set
y_train.shape
```

```
[ ]: (400, 1)
```

```
[ ]: # print the shapes of training and testing sets
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(400, 23)
(100, 23)
(400, 1)
(100, 1)
```

```
[ ]: # alternative way to split the dataset using train_test_split
from sklearn.model_selection import train_test_split
X = df[['OvertimeHoursPerMonth']]
y = df[['EducationLevel']]
X_train , X_test , y_train , y_test = train_test_split(X , y , test_size = 0.2 ,random_state = 42)
```

```
[ ]: # feature scaling using StandardScaler
sc = StandardScaler()
X_train_sc = sc.fit_transform(X_train)
X_test_sc = sc.fit_transform(X_test)

[ ]: # import DecisionTreeRegressor
from sklearn.tree import DecisionTreeRegressor

[ ]: # create Decision Tree Regressor model
model = DecisionTreeRegressor()

[ ]: # fit the model on training data
model.fit(X_train , y_train)

[ ]: DecisionTreeRegressor()

[ ]: # make predictions on the test set
y_pred = model.predict(X_test)

[ ]: # evaluate the model using R-squared metric
from sklearn.metrics import r2_score

[ ]: # calculate R-squared score
r2_score(y_test , y_pred)*100

[ ]: 1.3608422144420484

[ ]: # import LinearRegression
from sklearn.linear_model import LinearRegression

[ ]: # create Linear Regression model
model = LinearRegression()
model.fit(X_train_sc , y_train)

[ ]: LinearRegression()

[ ]: # make predictions on the test set
X_train_sc = sc.fit_transform(X_train)
X_test_sc = sc.fit_transform(X_test)

[ ]: # display first 10 rows of X_train
X_train[:10]

[ ]:      OvertimeHoursPerMonth
249                  5
433                  19
19                   2
```

```
322          34
332          16
56           16
301           5
229           4
331          10
132          22
```

```
[ ]: # check the dataframe info  
df.info()
```

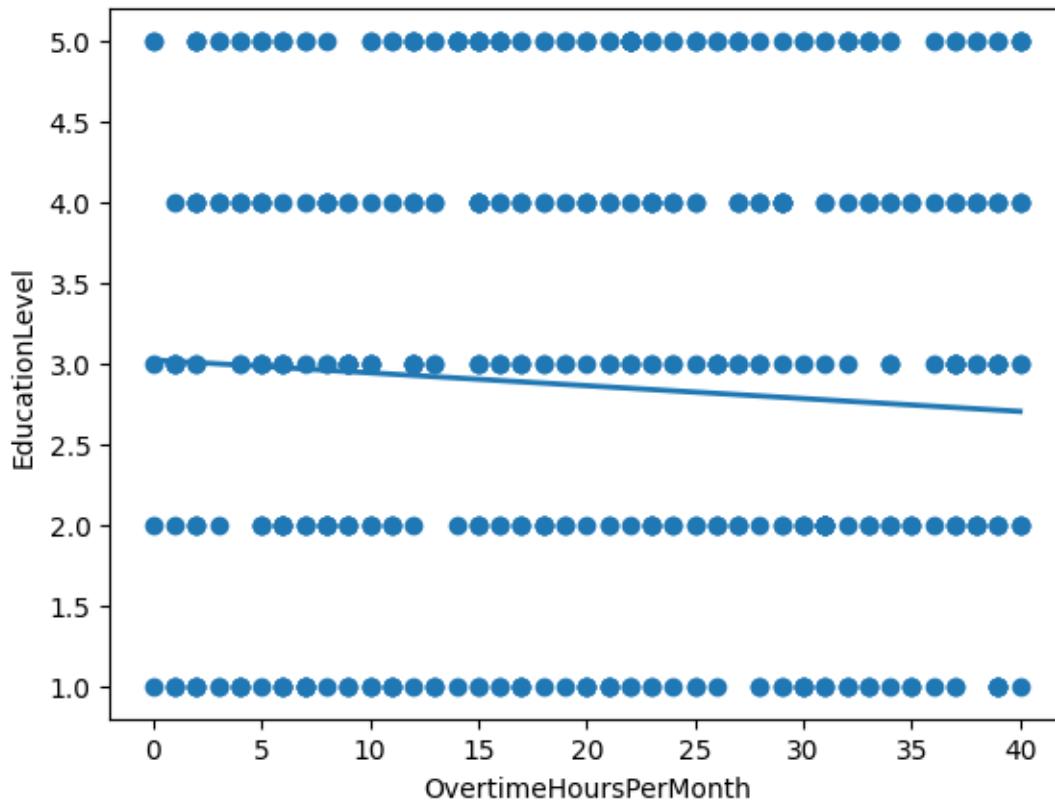
```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 500 entries, 0 to 499  
Data columns (total 24 columns):  
 #   Column            Non-Null Count  Dtype     
---  --     
 0   EmployeeID        500 non-null    int64    
 1   Name              500 non-null    object    
 2   Gender             500 non-null    object    
 3   Age               500 non-null    int64    
 4   Department         500 non-null    object    
 5   JobRole            500 non-null    object    
 6   EducationLevel     500 non-null    int64    
 7   JoiningDate        500 non-null    object    
 8   CountryCode        500 non-null    int64    
 9   Country            500 non-null    object    
 10  PhoneNumber        500 non-null    int64    
 11  MonthlySalary      500 non-null    int64    
 12  OvertimeHoursPerMonth  500 non-null    int64    
 13  LeavesTaken        500 non-null    int64    
 14  LastLeaveDate       500 non-null    object    
 15  LeaveDayName        500 non-null    object    
 16  ProjectsHandled     500 non-null    int64    
 17  TrainingHours       500 non-null    int64    
 18  CustomerSatisfaction  181 non-null    float64   
 19  LastPromotionYear    500 non-null    int64    
 20  YearsAtCompany       500 non-null    int64    
 21  WorkLifeBalanceScore  500 non-null    float64   
 22  PerformanceRating     500 non-null    int64    
 23  AttritionRisk        500 non-null    object    
dtypes: float64(2), int64(13), object(9)  
memory usage: 93.9+ KB
```

```
[ ]: # visualize the regression results  
plt.scatter(X.values , y.values , label = "Actual data")  
x_line = np.linspace(X.min().item(), X.max().item(), 100).reshape(-1,1)  
y_line = model.predict(x_line)
```

```

plt.plot(x_line , y_line , linewidth = 2 , label = "Regression line")
plt.xlabel("OvertimeHoursPerMonth")
plt.ylabel("EducationLevel")
plt.show()

```



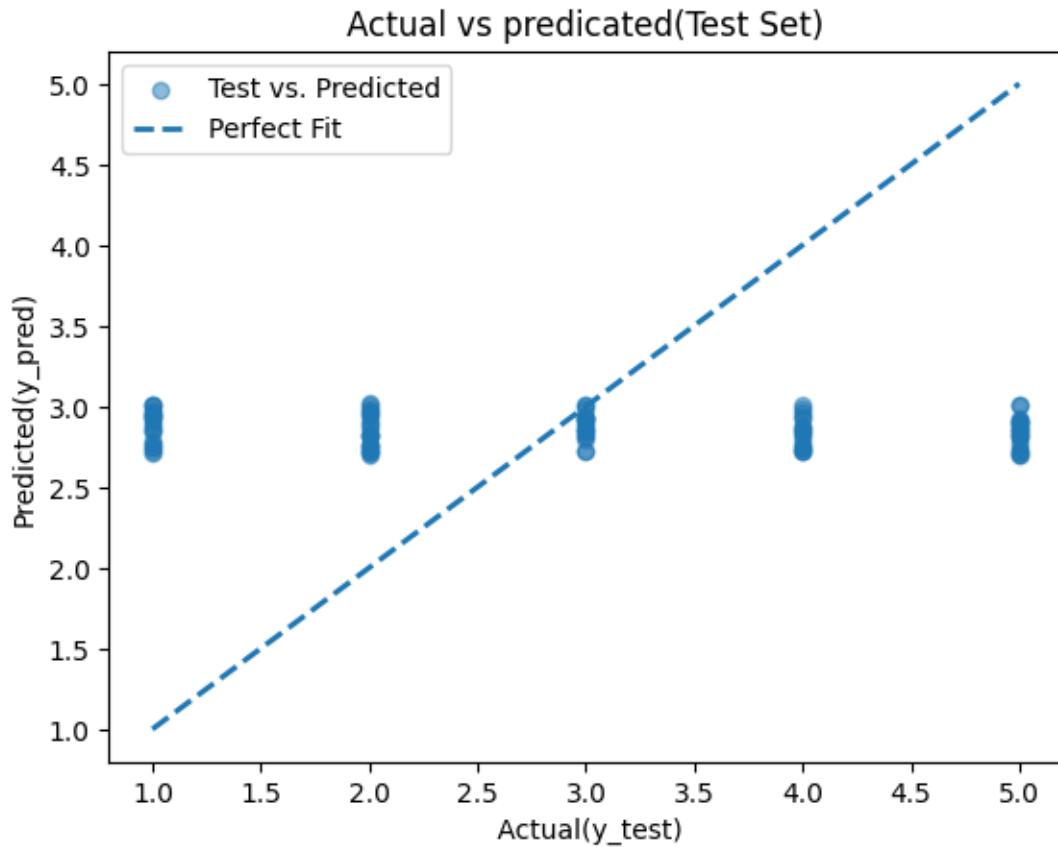
This plot shows that Education Level does not meaningfully change with Overtime Hours. The points are spread evenly across all education levels, and the trend line is almost flat. This means there is no real relationship—employees with any education level work similar amounts of overtime.

```

[ ]: # visualize Actual vs Predicted for test set
y_pred_test = model.predict(X_test).ravel()
plt.scatter(y_test.values.ravel() , y_pred_test , alpha = 0.5 , label = "Test vs.
↪ Predicted")
mn, mx = y_test.values.min() , y_test.values.max()
plt.plot([mn , mx] , [mn , mx] , linestyle = "--" , linewidth = 2, label =
↪ "Perfect Fit")
plt.xlabel(" Actual(y_test)")
plt.ylabel("Predicted(y_pred)")
plt.title("Actual vs predicated(Test Set)")
plt.legend()
plt.show()

```

```
/usr/local/lib/python3.12/dist-packages/sklearn/utils/validation.py:2732:  
UserWarning: X has feature names, but LinearRegression was fitted without  
feature names  
    warnings.warn(
```



This plot shows that the model's predictions are almost the same value every time, no matter what the actual education level is. The points cluster around ~3.0 instead of following the perfect-fit line. This means the model cannot learn the relationship and simply predicts the average for all cases — indicating very poor predictive performance.