

EDA_and_Prediction_Model_on_Galton_Family_Heights_Dataset

By Hatem Elgenedy

November 8, 2025

```
[ ]: #This cell imports the main Python libraries used
# for data analysis and visualization - NumPy and Pandas for handling data,
# Matplotlib and Seaborn for plotting, and it also hides unnecessary warning
# messages to keep the output clean.

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[ ]: #This cell loads the Galton Family Heights dataset from the given file path
# into a pandas DataFrame named df
# then displays the first six rows to preview the data.
df = pd.read_csv('/Users/hatemelgenedy/Desktop/AI and Data Science Microsoft
course/Galton_Family_Heights.csv')
df.head(6)
```

```
[ ]:   Father_height  Mother_height  Child_height  gender
0            78.5          67.0         73.2      1
1            78.5          67.0         69.2      0
2            78.5          67.0         69.0      0
3            78.5          67.0         69.0      0
4            75.5          66.5         73.5      1
5            75.5          66.5         72.5      1
```

```
[ ]: #This cell displays a summary of the dataset's structure, including column
# names, data types, number of non-null values, and the total number of
# entries -
# helping to check for missing data or data type issues.
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 934 entries, 0 to 933
Data columns (total 4 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 #   Column           Non-Null Count  Dtype  
 ---  -- 
```

```
0 Father_height    934 non-null      float64
1 Mother_height    934 non-null      float64
2 Child_height     934 non-null      float64
3 gender           934 non-null      int64
dtypes: float64(3), int64(1)
memory usage: 29.3 KB
```

```
[ ]: #This cell shows the dimensions of the dataset - the number of rows and columns
      ↪- helping you understand how many records and features are in the DataFrame.
df.shape
```

```
[ ]: (934, 4)
```

```
[ ]: #This cell lists the data type of each column (e.g., float, int, object),
      ↪helping you identify which columns are numeric, categorical, or other types
      ↪for further analysis.
df.dtypes
```

```
[ ]: Father_height      float64
      Mother_height       float64
      Child_height        float64
      gender              int64
      dtype: object
```

```
[122]: #this cell identifies and separates columns in the dataset by their data types - 
      ↪numeric, categorical, boolean, and datetime - then prints each list.
      #It helps you understand what kinds of data working with for analysis or
      ↪preprocessing.
```

```
num_cols = df.select_dtypes(include="number").columns
cat_cols = df.select_dtypes(include=["object", "category"]).columns
bool_cols = df.select_dtypes(include="bool").columns
dt_cols   = df.select_dtypes(include=["datetime", "datetimetz"]).columns

print("Numeric:", list(num_cols))
print("Categorical:", list(cat_cols))
print("Boolean:", list(bool_cols))
print("Datetime:", list(dt_cols))
```

```
Numeric: ['Father_height', 'Mother_height', 'Child_height', 'gender',
'Midparent_height']
Categorical: []
Boolean: []
Datetime: []
```

```
[ ]: #This cell installs the ydata_profiling library,
      # which is used to automatically generate a detailed exploratory data analysis
      ↪(EDA) report -
      # summarizing data distributions, correlations, and missing values in one step.
```

```
!pip install ydata_profiling
```

```
Requirement already satisfied: ydata_profiling in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (4.17.0)
Requirement already satisfied: scipy<1.16,>=1.4.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (1.13.1)
Requirement already satisfied: pandas!=1.4.0,<3.0,>1.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (2.3.1)
Requirement already satisfied: matplotlib<=3.10,>=3.5 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (3.10.0)
Requirement already satisfied: pydantic>=2 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (2.11.7)
Requirement already satisfied: PyYAML<6.1,>=5.0.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (6.0.2)
Requirement already satisfied: jinja2<3.2,>=2.11.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (3.1.6)
Requirement already satisfied: visions<0.8.2,>=0.7.5 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from visions[type_image_path]<0.8.2,>=0.7.5->ydata_profiling) (0.8.1)
Requirement already satisfied: numpy<2.2,>=1.16.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (1.26.4)
Requirement already satisfied: minify-html>=0.15.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (0.18.1)
Requirement already satisfied: filetype>=1.0.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (1.2.0)
Requirement already satisfied: phik<0.13,>=0.11.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (0.12.5)
Requirement already satisfied: requests<3,>=2.24.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (2.32.4)
Requirement already satisfied: tqdm<5,>=4.48.2 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (4.67.1)
Requirement already satisfied: seaborn<0.14,>=0.10.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (0.13.2)
Requirement already satisfied: multimethod<2,>=1.4 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata_profiling) (1.12)
Requirement already satisfied: statsmodels<1,>=0.13.2 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
```

```
ydata_profiling) (0.14.5)
Requirement already satisfied: typeguard<5,>=3 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata_profiling) (4.4.4)
Requirement already satisfied: imagehash==4.3.1 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata_profiling) (4.3.1)
Requirement already satisfied: wordcloud>=1.9.3 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata_profiling) (1.9.4)
Requirement already satisfied: dacite>=1.8 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata_profiling) (1.9.2)
Requirement already satisfied: numba<=0.61,>=0.56.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
ydata_profiling) (0.61.0)
Requirement already satisfied: PyWavelets in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from imagehash==4.3.1->ydata_profiling)
(1.9.0)
Requirement already satisfied: pillow in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from imagehash==4.3.1->ydata_profiling)
(11.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from jinja2<3.2,>=2.11.1->ydata_profiling)
(3.0.2)
Requirement already satisfied: contourpy>=1.0.1 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata_profiling)
(1.3.1)
Requirement already satisfied: cycler>=0.10 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata_profiling)
(0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
matplotlib<=3.10,>=3.5->ydata_profiling) (4.55.3)
Requirement already satisfied: kiwisolver>=1.3.1 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
matplotlib<=3.10,>=3.5->ydata_profiling) (1.4.8)
Requirement already satisfied: packaging>=20.0 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata_profiling)
(24.2)
Requirement already satisfied: pyparsing>=2.3.1 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata_profiling)
(3.2.0)
Requirement already satisfied: python-dateutil>=2.7 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
matplotlib<=3.10,>=3.5->ydata_profiling) (2.9.0.post0)
Requirement already satisfied: llvmlite<0.45,>=0.44.0dev0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
numba<=0.61,>=0.56.0->ydata_profiling) (0.44.0)
Requirement already satisfied: pytz>=2020.1 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from pandas!=1.4.0,<3.0,>1.1->ydata_profiling)
(2025.2)
```

```
Requirement already satisfied: tzdata>=2022.7 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from pandas!=1.4.0,<3.0,>1.1->ydata_profiling) (2025.2)
Requirement already satisfied: joblib>=0.14.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from phik<0.13,>=0.11.1->ydata_profiling) (1.5.1)
Requirement already satisfied: charset_normalizer<4,>=2 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from requests<3,>=2.24.0->ydata_profiling) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from requests<3,>=2.24.0->ydata_profiling) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from requests<3,>=2.24.0->ydata_profiling) (2.5.0)
Requirement already satisfied: certifi>=2017.4.17 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from requests<3,>=2.24.0->ydata_profiling) (2025.8.3)
Requirement already satisfied: patsy>=0.5.6 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from statsmodels<1,>=0.13.2->ydata_profiling) (1.0.2)
Requirement already satisfied: typing_extensions>=4.14.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from typeguard<5,>=3->ydata_profiling) (4.15.0)
Requirement already satisfied: attrs>=19.3.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2,>=0.7.5->ydata_profiling) (24.3.0)
Requirement already satisfied: networkx>=2.4 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2,>=0.7.5->ydata_profiling) (3.4.2)
Requirement already satisfied: puremagic in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2,>=0.7.5->ydata_profiling) (1.30)
Requirement already satisfied: annotated-types>=0.6.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from pydantic>=2->ydata_profiling) (0.6.0)
Requirement already satisfied: pydantic-core==2.33.2 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from pydantic>=2->ydata_profiling) (2.33.2)
Requirement already satisfied: typing-inspection>=0.4.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from pydantic>=2->ydata_profiling) (0.4.0)
Requirement already satisfied: six>=1.5 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from python-dateutil>=2.7->matplotlib<=3.10,>=3.5->ydata_profiling) (1.17.0)
```

```
[ ]: #This cell imports the ProfileReport class from the ydata_profiling library,  
#which allows you to create an automated EDA report for your dataset.  
from ydata_profiling import ProfileReport
```

```
[ ]: #This cell reloads the Galton Family Heights dataset into the DataFrame df and  
#displays the first six rows to confirm that the data was loaded correctly.  
df = pd.read_csv('/Users/hatemelgenedy/Desktop/AI and Data Science Microsoft  
course/Galton_Family_Heights.csv')  
df.head(6)
```

```
[ ]:   Father_height  Mother_height  Child_height  gender  
0            78.5          67.0        73.2      1  
1            78.5          67.0        69.2      0  
2            78.5          67.0        69.0      0  
3            78.5          67.0        69.0      0  
4            75.5          66.5        73.5      1  
5            75.5          66.5        72.5      1
```

```
[ ]: #this cell generates an automated EDA report for the dataset using ProfileReport  
#and saves it as an HTML file named "Galton Family EDA.html" -  
#giving a detailed summary of data patterns, distributions, and correlations.  
report = ProfileReport(df , title = 'Galton Family data analysis report')  
report.to_file(output_file = 'Galton Family EDA.html')
```

```
Summarize dataset:  0% | 0/5 [00:00<?, ?it/s]  
100% | 4/4 [00:00<00:00, 25003.30it/s]  
Generate report structure:  0% | 0/1 [00:00<?, ?it/s]  
Render HTML:  0% | 0/1 [00:00<?, ?it/s]  
Export report to file:  0% | 0/1 [00:00<?, ?it/s]
```

#The dataset represents Father height , Mother height , Child height , Gender

```
[ ]: #This cell checks for missing (null) values in each column of the dataset and  
#displays how many are present -  
# helping to identify if any data cleaning is needed.  
df.isnull().sum()
```

```
[ ]: Father_height    0  
Mother_height     0  
Child_height      0  
gender            0  
dtype: int64
```

```
[ ]: #This cell checks for duplicate rows in the dataset and returns a boolean  
#Series indicating which rows are duplicates -  
df.duplicated()
```

```
[ ]: 0      False
    1      False
    2      False
    3      True
    4      False
    ...
929     False
930     False
931     False
932     False
933     False
Length: 934, dtype: bool
```

```
[ ]: #This cell shows the number of rows and columns in the dataset - helping to
      ↪confirm the dataset's overall size.
df.shape
```

```
[ ]: (934, 4)
```

```
[ ]: #This cell checks the dataset for duplicate rows and prints the total number
      ↪found -
#helping to identify repeated data that might need removal.
duplicates = df . duplicated().sum()
print(f"Number of duplicate rows : {duplicates}")
```

```
Number of duplicate rows : 182
```

```
[ ]: #This cell creates a new column called Midparent_height, which calculates the
      ↪average of the father's and mother's heights -
#representing the parents' combined average height.
df["Midparent_height"] = (df["Father_height"] + df["Mother_height"]) / 2
```

```
[ ]: #This cell calculates the correlation matrix between all numeric columns in the
      ↪dataset and prints it - showing:
#how strongly each pair of variables (like parent and child heights) are
      ↪related.
corr = df.corr(numeric_only=True)
print("\nCorrelation Matrix:\n", corr)
```


Correlation Matrix:

	Father_height	Mother_height	Child_height	gender	\
Father_height	1.000000	0.060366	0.266039	-0.024762	
Mother_height	0.060366	1.000000	0.201322	-0.025216	
Child_height	0.266039	0.201322	1.000000	0.716709	
gender	-0.024762	-0.025216	0.716709	1.000000	
Midparent_height	0.752750	0.702546	0.322443	-0.034284	

	Midparent_height
Father_height	0.752750
Mother_height	0.702546
Child_height	0.322443
gender	-0.034284
Midparent_height	1.000000

```
[ ]: #This cell computes the correlation matrix for all numeric columns in the dataset -
#measuring the strength and direction of relationships between numerical variables.
corr = df.corr(numeric_only=True)
```


```
[ ]: #This cell visualizes the correlation matrix as a heatmap, where colors represent how strongly variables are related -
#with coolwarm showing positive (red) and negative (blue) correlations.
 #%matplotlib inline ensures the plot appears directly in the notebook.
%matplotlib inline
plt.figure(figsize=(6,4))
sns.heatmap(corr, annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```




```
[ ]: # X axis represents the height of the father,Mother , Child
# Y axis. represents how many fathers,Mothers , Childrens in this heights
#This cell creates histograms for the heights of fathers, mothers, and children
#to visualize how each is distributed.
#The kde=True adds a smooth density curve, and plt.tight_layout() ensures the
#plots don't overlap for a cleaner display.
fig = plt.figure(figsize=(14, 4))
axes = fig.subplots(1, 4)
sns.histplot(df["Father_height"], kde=True, ax=axes[0])
plt.title("Father Height")

sns.histplot(df["Mother_height"], kde=True, ax=axes[1])
plt.title("Mother Height")


sns.histplot(df["Child_height"], kde=True, ax=axes[2])
plt.title("Child Height")
plt.tight_layout()
plt.show()
```



```
[ ]: #Outliers in Father,Mother,Child height
#This cell creates boxplots for father, mother, and child heights to visually
#identify outliers and compare height distributions -
#The boxes show data spread,
#while points outside indicate possible outliers.
fig = plt.figure(figsize=(8, 4))
ax = fig.subplots()
sns.boxplot(data=df[["Father_height", "Mother_height", "Child_height"]], ax=ax)
plt.title("Boxplots: Outlier Check")
plt.ylabel("Height (in)")
plt.show()
```



```
[ ]: #Gender count
#This cell creates a count plot showing how many records belong to each gender ↴
#(e.g., male vs. female).
#It helps visualize the gender distribution in the dataset.
fig = plt.figure(figsize=(8, 4))
ax = fig.subplots()
sns.countplot(x="gender", data=df, ax=ax)
ax.set_title("Gender-wise Count")
ax.set_xlabel("Gender")
ax.set_ylabel("Count")
plt.show()
```




```
[ ]: #This cell imports key tools from scikit-learn:  
#train_test_split to divide data into training and testing sets,  
#LinearRegression and RandomForestRegressor to build prediction models,  
#r2_score and mean_absolute_error to evaluate model accuracy and error.  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression  
from sklearn.ensemble import RandomForestRegressor  
from sklearn.metrics import r2_score,mean_absolute_error
```

```
[ ]: #This cell displays the first few rows of the dataset to confirm that the data  
↳ is loaded correctly and to get an initial look at the structure and values  
↳ in the DataFrame.  
df.head()
```

	Father_height	Mother_height	Child_height	gender	Midparent_height
0	78.5	67.0	73.2	1	72.75
1	78.5	67.0	69.2	0	72.75
2	78.5	67.0	69.0	0	72.75
3	78.5	67.0	69.0	0	72.75
4	75.5	66.5	73.5	1	71.00

```
[ ]: #This scatter plot shows the relationship between father height and child  
↳ height - each point represents one family.  
#It helps visualize how a child's height tends to increase as the father's  
↳ height increases, showing a positive correlation.  
plt.figure(figsize=(10,5))
```

```
plt.show()
```




<Figure size 1000x500 with 0 Axes>

<Figure size 1000x500 with 0 Axes>

```
[ ]: #This cell selects only the numeric columns (father, mother, child heights, and gender) and computes their correlation matrix, showing how strongly these variables are related to each other.  
num_cols = ["Father_height", "Mother_height", "Child_height", "gender"]  
corr = df[num_cols].corr(numeric_only=True)
```

```
[ ]: #This cell creates a heatmap of the numeric columns' correlation values - visually showing how closely each variable (father, mother, child heights, and gender) is related, using color intensity to represent the correlation strength.  
plt.figure(figsize=(5,4))  
sns.heatmap(corr, annot=True, cmap="coolwarm", fmt=".2f")  
plt.title("Correlation Heatmap")  
plt.tight_layout()  
plt.show()
```



```
[ ]: #average child height by gender
#This cell groups the data by gender and calculates the average child height
#↳ for each group,
#rounding the results to two decimal places - helping compare the mean height
#↳ between males and females.
avg_by_gender = df.groupby("gender")["Child_height"].mean().round(2)
print("\nAverage Child Height by Gender:\n", avg_by_gender)
```

Average Child Height by Gender:

gender	Child_height
0	64.10
1	69.23

Name: Child_height, dtype: float64

```
[ ]: #Model Building - Child Height Prediction
#This cell defines the features (X) - father's height, mother's height, and
#↳ gender -
#and the target variable (y), which is the child's height, preparing the data
#↳ for model training.
X = df[["Father_height", "Mother_height", "gender"]]
y = df["Child_height"]
```

```
[ ]: #This cell splits the dataset into training and testing sets, with 80% of the data used for training and 20% for testing.
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42)

[ ]: #This cell initializes a Linear Regression model, which will be used to predict child heights based on the features defined in X.
linreg = LinearRegression()

[ ]: # This cell initializes a Random Forest Regressor model, which will be used to predict child heights based on the features defined in X.
rf = RandomForestRegressor(random_state=42, n_estimators=300)

[ ]: #This cell fits the Linear Regression model to the training data (X_train and y_train) and also fits the Random Forest Regressor to the same training data, allowing both models to learn the relationships between the features and the target variable.
linreg.fit(X_train, y_train)
rf.fit(X_train, y_train)

[ ]: RandomForestRegressor(n_estimators=300, random_state=42)

[ ]: #This cell uses the trained Linear Regression and Random Forest models to make predictions on the test set (X_test), storing the predicted child heights in y_pred_lr and y_pred_rf respectively.
y_pred_lr = linreg.predict(X_test)
y_pred_rf = rf.predict(X_test)

[ ]: #This cell evaluates the performance of the Linear Regression model by calculating the R-squared score and Mean Absolute Error (MAE) between the actual child heights (y_test) and the predicted heights (y_pred_lr).
r2_lr = r2_score(y_test, y_pred_lr)
mae_lr = mean_absolute_error(y_test, y_pred_lr)

[ ]: #This cell evaluates the performance of the Random Forest Regressor model by calculating the R-squared score and Mean Absolute Error (MAE) between the actual child heights (y_test) and the predicted heights (y_pred_rf).
r2_rf = r2_score(y_test, y_pred_rf)
mae_rf = mean_absolute_error(y_test, y_pred_rf)

[ ]: #This cell prints the R-squared and MAE values for both the Linear Regression and Random Forest models, allowing you to compare their performance in predicting child heights.
print(f"Linear Regression R2: {r2_lr:.3f}, MAE: {mae_lr:.3f}")
print(f"Random Forest R2: {r2_rf:.3f}, MAE: {mae_rf:.3f}")
```

Linear Regression R2: 0.613, MAE: 1.687

Random Forest R2: 0.590, MAE: 1.670

```
[ ]: #This cell compares the R-squared scores of both models and prints which one
    ↪performed better in terms of R2, indicating which model explains more
    ↪variance in the child height data.
best_model_name = "Random Forest" if r2_rf >= r2_lr else "Linear Regression"
print(f"\nBetter model (by R2): {best_model_name}")
```

Better model (by R2): Linear Regression

#The x-axis is the actual child height from test data. #The y-axis is the predicted height from model. #The blue line is the perfect prediction line — where Actual = Predicted.


#The closer the dots are to the diagonal line, the better the predictions. #Dots above the line → model predicted too high. #Dots below the line → model predicted too low.

#The points are mostly near the diagonal, meaning the Linear Regression model predicts well.

#Clear positive trend: as actual height increases, predicted height also increases.

```
[ ]: #Model Interpretation & Visualization
#Plot actual vs predicted (use best model)
#This cell creates a scatter plot comparing the actual child heights (y_test) ↪
    ↪to the predicted heights from the best-performing model (either Random ↪
    ↪Forest or Linear Regression).
if best_model_name == "Random Forest":
    y_pred_best = y_pred_rf
else:
    y_pred_best = y_pred_lr

plt.figure(figsize=(5,5))
plt.scatter(y_test, y_pred_best, s=12)
# diagonal reference line
mn, mx = min(y_test.min(), y_pred_best.min()), max(y_test.max(), y_pred_best.
    ↪max())
plt.plot([mn, mx], [mn, mx])
plt.title(f"Actual vs Predicted (Best: {best_model_name})")
plt.xlabel("Actual Child Height (in)")
plt.ylabel("Predicted Child Height (in)")
plt.tight_layout()
plt.show()
```



#The x-axis is the index of each test record. #The y-axis is the residual (error) = Actual – Predicted.

#If residual = 0 → perfect prediction.

#Dots above 0 → model under-predicted (predicted too low). #Dots below 0 → model over-predicted (predicted too high).

#Most residuals are centered around 0.

#No clear pattern (they spread randomly above and below the line), which means the model doesn't have systematic bias.

#A few points are farther away from 0 — those are larger errors or outliers.

[]: *#This cell calculates the residuals (the difference between actual and predicted values) for the best model and creates a scatter plot of these residuals to check for patterns or biases in the predictions.*


```
residuals = y_test - y_pred_best
```

```
plt.figure(figsize=(5,4))
```

```

plt.scatter(range(len(residuals)), residuals, s=12)
plt.axhline(0)
plt.title("Residual Errors (Actual - Predicted)")
plt.xlabel("Test Sample Index")
plt.ylabel("Residual")
plt.tight_layout()
plt.show()


```



#Residual = Actual – Predicted #X-axis → the size of those residuals (how far off each prediction was)
#Y-axis → how many data points had that size of error
#Bars near 0 → predictions close to actual values.
#Bars farther from 0 → larger errors.
#Most residuals cluster between about -3 and +3 → predictions are close to real heights.
#The peak is near 0, showing the model is centered correctly (no consistent over/under-prediction).
#The shape is almost normal (bell-shaped), which supports Linear Regression assumptions are reasonable.
#A few bars on the far left/right (like -10 or +7) are outliers — these are just a few cases where the model missed badly.
#This histogram shows that most prediction errors are small and evenly spread around zero.

#The model performs well overall, with only a few large outliers.

```
[ ]: #This cell creates a histogram of the residuals to visualize their distribution, helping to check if they are normally distributed around zero, which is an assumption of many regression models.  
plt.figure(figsize=(5,4))  
plt.hist(residuals, bins=20)  
plt.title("Residuals Histogram")  
plt.xlabel("Residual")  
plt.ylabel("Count")  
plt.tight_layout()  
plt.show()
```



```
[ ]: #This cell creates a new DataFrame called comparison that contains the actual child heights (y_test) and the predicted heights (y_pred_best) from the best model, then prints the first 10 rows of this comparison to show how well the predictions match the actual values.  
#comparing actual vs predicted  
comparison = pd.DataFrame({  
    "Actual": y_test.values,  
    "Predicted": y_pred_best  
}).reset_index(drop=True)
```

```
print("\nFirst 10 - Actual vs Predicted:")
print(comparison.head(10).round(2))
```

First 10 - Actual vs Predicted:

	Actual	Predicted
0	65.5	64.32
1	72.7	72.21
2	70.0	69.72
3	69.5	68.50
4	64.0	67.29
5	63.0	64.85
6	64.5	64.71
7	65.5	65.44
8	63.5	63.52
9	74.0	69.81